

# Internetworking: Packet Switching & Performance Metrics

EE4204: Computer Networks

Mehul Motani

motani@nus.edu.sg


**Note:** Some slides & graphics adapted from “Computer Networking” by Kurose & Ross

EE4204 Lecture Notes

Computer Networks


## A closer look at network structure

- **network edge:**  
applications and hosts
- **access networks, physical media:**  
wired, wireless communication links
- **network core:**
  - interconnected routers
  - network of networks




# The Network Core

- mesh of interconnected routers
- **fundamental question:**  
how is data transferred through net?
  - **circuit switching:**  
dedicated circuit per call: telephone net
  - **packet-switching:**  
data sent thru net in discrete “chunks”



## Network Core: Circuit Switching

- End-end resources reserved for “call”**
- link bandwidth, switch capacity
  - dedicated resources: no sharing
  - circuit-like (guaranteed) performance
  - call setup required




# Circuit Switching: FDM and TDM


FDM

Example:

4 users



TDM



## Network Core: Packet Switching

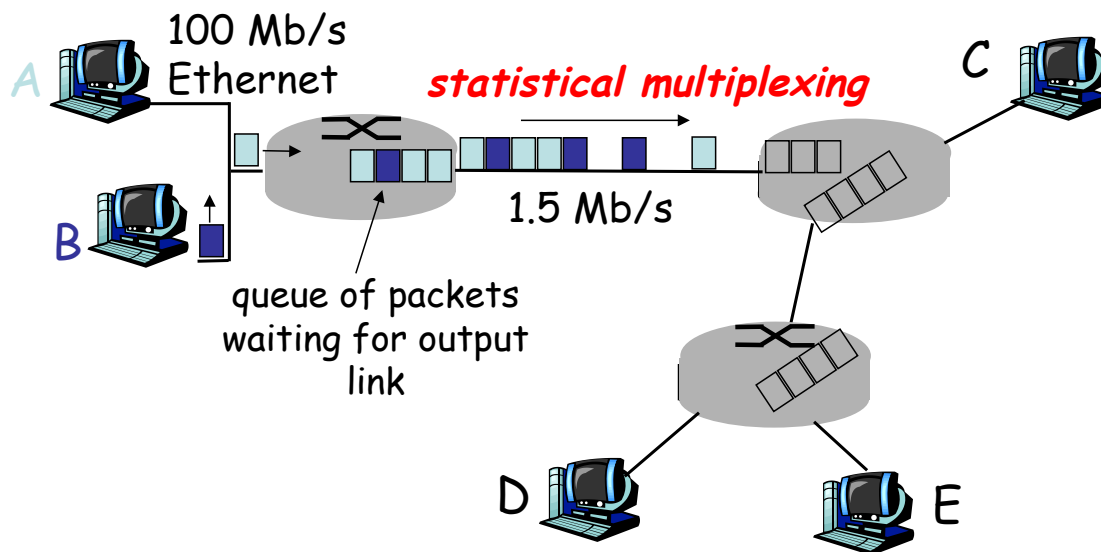
each end-end data stream  
divided into *packets*

- user A, B packets *share* network resources
- each packet uses full link bandwidth
- resources used *as needed*

resource contention:

- aggregate resource demand can exceed amount available
- congestion: packets queue, wait for link use
- store and forward: packets move one hop at a time
  - Node receives complete packet before forwarding

# Packet Switching: Statistical Multiplexing




Sequence of A & B packets does not have fixed pattern,  
bandwidth shared on demand ➔ **statistical multiplexing**.  
TDM: each host gets same slot in revolving TDM frame.

## Packet switching versus circuit switching

**Packet switching allows more users to use network!**

- 1 Mb/s link
- each user:
  - 100 kb/s when “active”
  - active 10% of time
- **circuit-switching:**
  - 10 users
- **packet switching:**
  - with 35 users, probability > 10 active at same time is less than .0005



Q: how did we get value 0.0005?

# Packet switching versus circuit switching


## Is packet switching a “slam dunk winner?”

- great for bursty data
  - resource sharing
  - simpler, no call setup
- **excessive congestion**: packet delay and loss
  - protocols needed for reliable data transfer, congestion control
- **Q: How to provide circuit-like behavior?**
  - bandwidth guarantees needed for audio/video apps
  - still an unsolved problem

# How do delay and loss occur?

## Packets queue in router buffers in network core

- packet arrival rate to link exceeds output link capacity
- packets queue, wait for turn



# Delay in packet-switched networks

## 1. nodal processing:

- check bit errors
- determine output link

## 2. queueing


- time waiting at output link for transmission
- depends on congestion level of router

## 3. Transmission delay:

- $R$  = link bandwidth (bps)
- $L$  = packet length (bits)
- time to send bits into link =  $L/R$

## 4. Propagation delay:

- $d$  = length of physical link
- $s$  = propagation speed in medium ( $\sim 2 \times 10^8$  m/sec)
- propagation delay =  $d/s$



**Note:**  $s$  and  $R$  are very different quantities!


# Nodal delay

$$d_{\text{nodal}} = d_{\text{proc}} + d_{\text{queue}} + d_{\text{trans}} + d_{\text{prop}}$$

- $d_{\text{proc}}$  = processing delay
  - typically a few microseconds or less
- $d_{\text{queue}}$  = queuing delay
  - depends on congestion
- $d_{\text{trans}}$  = transmission delay
  - $= L/R$ , significant for low-speed links
- $d_{\text{prop}}$  = propagation delay
  - a few microsecs to hundreds of msecs


# Packet loss

- queue (aka buffer) preceding link in buffer has finite capacity
- packet arriving to full queue dropped (aka lost)
- lost packet may be retransmitted by previous node, by source end system, or not at all



## “Real” Internet delays and routes

- **Ping program**: hello message to remote host, gives delay and loss rate
- **Traceroute program**: provides delay measurement from source to router along end-end Internet path towards destination. For all  $i$ :
  - sends three probes that will reach router  $i$  on path towards destination
  - router  $i$  will return packets to sender
  - sender times interval between transmission and reply.





# “Real” Internet delays and routes

```
[frogg:~ motani$ ping www.yahoo.com
PING atsv2-fp-shed.wg1.b.yahoo.com (106.10.250.10): 56 data bytes
64 bytes from 106.10.250.10: icmp_seq=0 ttl=55 time=6.971 ms
64 bytes from 106.10.250.10: icmp_seq=1 ttl=55 time=9.234 ms
64 bytes from 106.10.250.10: icmp_seq=2 ttl=55 time=9.456 ms
64 bytes from 106.10.250.10: icmp_seq=3 ttl=55 time=16.526 ms
64 bytes from 106.10.250.10: icmp_seq=4 ttl=55 time=10.097 ms
64 bytes from 106.10.250.10: icmp_seq=5 ttl=55 time=9.866 ms
64 bytes from 106.10.250.10: icmp_seq=6 ttl=55 time=6.971 ms
64 bytes from 106.10.250.10: icmp_seq=7 ttl=55 time=9.301 ms
64 bytes from 106.10.250.10: icmp_seq=8 ttl=55 time=18.473 ms
64 bytes from 106.10.250.10: icmp_seq=9 ttl=55 time=49.639 ms
64 bytes from 106.10.250.10: icmp_seq=10 ttl=55 time=8.990 ms
64 bytes from 106.10.250.10: icmp_seq=11 ttl=55 time=9.174 ms
64 bytes from 106.10.250.10: icmp_seq=12 ttl=55 time=9.903 ms
64 bytes from 106.10.250.10: icmp_seq=13 ttl=55 time=27.054 ms
64 bytes from 106.10.250.10: icmp_seq=14 ttl=55 time=30.060 ms
64 bytes from 106.10.250.10: icmp_seq=15 ttl=55 time=9.538 ms
64 bytes from 106.10.250.10: icmp_seq=16 ttl=55 time=9.555 ms
64 bytes from 106.10.250.10: icmp_seq=17 ttl=55 time=9.547 ms
64 bytes from 106.10.250.10: icmp_seq=18 ttl=55 time=16.386 ms
^C
--- atsv2-fp-shed.wg1.b.yahoo.com ping statistics ---
19 packets transmitted, 19 packets received, 0.0% packet loss
round-trip min/avg/max/stddev = 6.971/14.565/49.639/10.376 ms
frogg:~ motani$
```

# “Real” Internet delays and routes

**traceroute:** gaia.cs.umass.edu to www.eurecom.fr

Three delay measurements from  
gaia.cs.umass.edu to cs-gw.cs.umass.edu

```

1 cs-gw (128.119.240.254) 1 ms 1 ms 2 ms
2 border1-rt-fa5-1-0.gw.umass.edu (128.119.3.145) 1 ms 1 ms 2 ms
3 cht-vbns.gw.umass.edu (128.119.3.130) 6 ms 5 ms 5 ms
4 jn1-at1-0-0-19.wor.vbns.net (204.147.132.129) 16 ms 11 ms 13 ms
5 jn1-so7-0-0-0.wae.vbns.net (204.147.136.136) 21 ms 18 ms 18 ms
6 abilene-vbns.abilene.ucaid.edu (198.32.11.9) 22 ms 18 ms 22 ms
7 nycm-wash.abilene.ucaid.edu (198.32.8.46) 22 ms 22 ms 22 ms
8 62.40.103.253 (62.40.103.253) 104 ms 109 ms 106 ms
9 de2-1.de1.de.geant.net (62.40.96.129) 109 ms 102 ms 104 ms
10 de.fr1.fr.geant.net (62.40.96.50) 113 ms 121 ms 114 ms
11 renater-gw.fr1.fr.geant.net (62.40.103.54) 112 ms 114 ms 112 ms
12 nio-n2.cssi.renater.fr (193.51.206.13) 111 ms 114 ms 116 ms
13 nice.cssi.renater.fr (195.220.98.102) 123 ms 125 ms 124 ms
14 r3t2-nice.cssi.renater.fr (195.220.98.110) 126 ms 126 ms 124 ms
15 eurecom-valbonne.r3t2.ft.net (193.48.50.54) 135 ms 128 ms 133 ms
16 194.214.211.25 (194.214.211.25) 126 ms 128 ms 126 ms
17 * * *
18 * * *
19 fantasia.eurecom.fr (193.55.113.142) 132 ms 128 ms 136 ms
  
```


trans-oceanic link

\* means no response (probe lost, router not replying)



# Throughput

- **throughput**: rate (bits/time) at which bits transferred between sender/receiver
- **instantaneous**: rate at given point in time
- **average**: rate over long(er) period of time




## **bottleneck link**

link on end-end path that constrains end-end throughput

# Throughput: Internet scenario

- per-connection end-end throughput:  $\min(R_c, R_s, R/10)$
- in practice:  $R_c$  or  $R_s$  is often bottleneck




10 connections (fairly) share backbone bottleneck link  $R$  bits/sec

# Internet History

<http://edgalaxy.com/history/2010/9/16/the-history-of-the-internet.html>

## 1961-1972: Early packet-switching principles

- 1961: Kleinrock - queueing theory shows effectiveness of packet-switching
- 1964: Baran - packet-switching in military nets
- 1967: ARPAnet conceived by Advanced Research Projects Agency
- 1969: first ARPAnet node operational
- 1972:
  - ARPAnet public demonstration
  - NCP (Network Control Protocol) first host-host protocol
  - first e-mail program
  - ARPAnet has 15 nodes



THE ARPA NETWORK

# Internet History

## 1972-1980: Internetworking, proprietary networks

- 1970: ALOHAnet satellite network in Hawaii
- 1974: Cerf and Kahn - architecture for interconnecting networks
- 1976: Ethernet at Xerox PARC
- late 70's: proprietary architectures: DECnet, SNA, XNA
- late 70's: switching fixed length packets (ATM precursor)
- 1979: ARPAnet has 200 nodes

### Cerf and Kahn's internetworking principles:

- minimalism, autonomy - no internal changes required to interconnect networks
- best effort service model
- stateless routers
- decentralized control

define today's Internet architecture

# Internet History

## *1980-1990: new protocols and new networks*

- 1983: deployment of TCP/IP
- 1982: smtp e-mail protocol defined
- 1983: DNS defined for name-to-IP-address translation
- 1985: ftp protocol defined
- 1988: TCP congestion control
- new national networks: Cset, BITnet, NSFnet, Minitel
- 100,000 hosts connected to confederation of networks

# Internet History

## *1990, 2000's: Internet explosion, web apps*

- Early 1990's: ARPAnet decommissioned
- 1991: NSF lifts restrictions on commercial use of NSFnet (decommissioned, 1995)
- early 1990s: Web
  - hypertext [Bush 1945, Nelson 1960's]
  - HTML, HTTP: Berners-Lee
  - 1994: Mosaic, later Netscape
  - late 1990's: commercialization of the Web
- Late 1990's – 2000's:
  - more killer apps: instant messaging, P2P file sharing
  - network security to forefront
  - est. 50 million host, 100 million+ users
  - backbone links running at Gbps

# Internet History

## *2010 & beyond: rebirth of the web*

- Over 1 billion hosts
- Voice, Video over IP
- P2P applications: BitTorrent (file sharing) Skype (VoIP), PPLive (video)
- Rich content & applications: Gaming, YouTube, Facebook, Twitter, Google
- Mobile wireless broadband
- Web 2.0, Internet 2.0
- Clean slate Internet – a complete redesign?